

Do Not Let English Judge the World: Protocol Sensitivity in Multilingual LLM-as-a-Judge Evaluation

Anonymous authors

Paper under double-blind review

Abstract

LLM-as-a-judge evaluation is increasingly used to score open-ended model outputs, but multilingual evaluation often hides a consequential implementation choice: which language the judge sees, which language the rubric uses, and whether non-English examples are translated into English before judging. We study this *protocol sensitivity* using SEAHORSE multilingual summarization ratings, source-grounded XLSum examples, and a WMT MQM translation-quality contrast. Across 580 main-experiment base examples and 2,560 judge responses, a gpt-4o-mini judge gives materially different scores under direct English-rubric judging, target-language rubrics, explicit English-pivot judging, and bilingual rubrics. The largest shifts occur on target-language form dimensions: Vietnamese grammar scores differ by 1.20 points on a 1–5 scale between target-language-rubric and English-pivot protocols. English pivoting is not a safe default: for Turkish comprehensibility it significantly reduces alignment with human labels, and follow-up experiments with gpt-4.1-mini, gpt-5.4-mini, and gpt-5.5 reproduce the key failure. A small evidence-preserving bilingual mitigation recovers the worst gpt-5.5 pivot failures. Source-grounded semantics and WMT MQM show smaller or more model-dependent effects, clarifying the boundary of the claim. Across 17 main-experiment task/language cells, the best protocol is not direct in 11 cells and explicit pivot is worst in 10. We argue that global evaluation papers should report judge protocol, per-language alignment, same-item protocol shifts, calibration behavior, parse rates, and token cost instead of a single aggregate judge score.

1 Introduction

LLM-as-a-judge evaluation has become an attractive default for open-ended generation: it is cheap, fast, repeatable at scale, and can be adapted to many tasks with a rubric prompt. Prior work shows that prompted LLM judges can correlate with human preferences and task-specific human ratings in chatbot evaluation, summarization, and machine translation (Zheng et al., 2023; Liu et al., 2023; Kocmi & Federmann, 2023). This has made LLM judging a common component of research leaderboards, dataset construction, model selection, and product evaluation.

The standard LLM-judge framing, however, is underspecified for global use. In an English-only benchmark, the evaluator usually sees instructions, rubric, candidate answer, and any reference material in the same language. In multilingual evaluation, these surfaces can be separated. A benchmark builder can ask an English prompt to judge a non-English answer directly; translate the candidate, source, or reference into English and judge the translation; translate the rubric into the target language; or provide bilingual instructions. These choices are often treated as engineering details, but they encode different assumptions about what evidence the judge should use.

The risk is clearest when the evaluation target is not purely semantic. English-pivot judging may help an English-dominant model process content, but the translation step can also remove exactly the evidence needed to judge grammar, fluency, repetition, morphology, reg-

ister, or whether the generated text is natural in the target language. A disfluent Vietnamese or Turkish summary can become a clean English sentence after translation; conversely, a good original can be damaged by a noisy pivot translation. Even when the target is semantic adequacy, the rubric language and translation pipeline may change the score distribution or the judge’s calibration. A multilingual benchmark score is therefore not just a model score; it is a score produced by a *judge protocol*.

This paper asks a practical question for global benchmark builders: *when a paper reports an LLM judge score for multilingual outputs, how sensitive is that score to the language protocol used to obtain it?* We do not propose a new judge model. Instead, we treat protocol choice as an experimental factor that should be measured and reported.

We study three settings with existing human labels. First, we use SEAHORSE candidate-quality labels for multilingual summaries, focusing on candidate-visible dimensions where target-language form evidence should matter. Second, we reconstruct source-grounded XLSum examples for the SEAHORSE main-ideas dimension to test a more semantic setting. Third, we use a small WMT MQM translation-quality contrast to ask whether machine-translation evaluation behaves differently when a reference is available. We compare four judge protocols: direct English-rubric judging, target-language-rubric judging, explicit English-pivot judging, and bilingual-rubric judging.

Our contributions are:

- We define *protocol sensitivity* for multilingual LLM-as-a-judge evaluation and operationalize it with same-item score shifts, by-language human alignment, language gaps, calibration diagnostics, parse rates, and token cost.
- We release a benchmark scaffold covering 580 main base examples across candidate-quality, source-grounded semantic, and WMT MQM translation-quality settings, plus stronger-judge, modern-judge, and mitigation experiments.
- We show that English pivoting is not a safe default: it significantly hurts Turkish comprehensibility in the main run, and the same failure mode appears under stronger and current frontier judges.
- We provide a Global Judge Reporting Card that makes multilingual judge protocols auditable rather than hidden prompt details.

Our main finding is not that one protocol always wins. The important result is that protocol choice can change both absolute scores and human alignment, and the direction depends on language, task, dimension, and judge model.

2 Related Work and Positioning

LLM-as-a-judge evaluation. Prompted LLM judges have been used for pairwise preference evaluation, rubric-based scoring, summarization evaluation, and machine-translation evaluation (Zheng et al., 2023; Liu et al., 2023; Kocmi & Federmann, 2023). This line of work demonstrates the practical value of LLM evaluation while also documenting sensitivity to prompt design, verbosity, position, and self-preference biases (Wang et al., 2023; Bavaresco et al., 2024). Our work studies a different source of sensitivity: the language protocol used to present multilingual evidence to the judge.

Multilingual and global evaluation. Multilingual benchmarks such as XTREME and XLSum expanded evaluation beyond English, but benchmark construction often relies on translation, parallel data, or language-specific preprocessing choices that affect what is measured (Hu et al., 2020; Hasan et al., 2021). SEAHORSE is especially useful for our study because it provides multilingual, multifaceted human labels for summarization quality (Clark et al., 2023). Recent work on multilingual judges finds that LLM judges can disagree across languages and that multilingual adjudication procedures can improve reliability (Fu & Liu, 2025; Fu & Lu, 2026). Our contribution is complementary: we evaluate low-cost reporting choices available to any benchmark builder, rather than training or designing a new judge.

94 **Human evaluation for summarization and MT.** Summarization and translation evalua-
 95 tion illustrate why protocol choice matters. Human evaluation studies distinguish semantic
 96 adequacy, factual consistency, fluency, grammar, and other quality dimensions (Fabbri et al.,
 97 2021; Maynez et al., 2020). MT evaluation has similarly moved toward expert error anno-
 98 tation and MQM-style judgments, where references, source context, and target-language
 99 form can all matter (Freitag et al., 2021; 2022). These settings let us ask whether protocol
 100 sensitivity is strongest when the label depends on target-language form and weaker when a
 101 reference constrains the judgment.

102 3 Experimental Setup

103 3.1 Research Questions

104 We organize the experiments around four questions. **RQ1:** How much does the judge
 105 protocol change same-item scores and human alignment? **RQ2:** Are protocol effects larger
 106 for target-language form dimensions than for source-grounded semantic or reference-based
 107 translation-quality judgments? **RQ3:** Do stronger judges or threshold calibration remove
 108 the failure mode? **RQ4:** What should future multilingual LLM-judge papers report so that
 109 their conclusions are auditable?

110 3.2 Datasets and Sampling

111 **Candidate-quality setting.** We sample from the SEAHORSE test split (Clark et al., 2023).
 112 The main run uses English (en-US), Spanish (es-ES), Turkish (tr), and Vietnamese (vi). We
 113 evaluate two candidate-visible dimensions: comprehensibility and grammar. For each
 114 language and dimension, we sample 50 balanced examples, with 25 positive and 25 negative
 115 human labels. This yields 400 base examples and 1,600 judge responses across four protocols.

116 **Source-grounded semantic setting.** SEAHORSE raw TSV files do not include source doc-
 117 uments. To recover source-grounded examples, we match SEAHORSE reference summaries
 118 to raw XLSum test summaries (Hasan et al., 2021). Matching recovers 151/151 Spanish,
 119 243/244 Turkish, and 252/252 Vietnamese reference rows. We then sample 30 balanced
 120 main-ideas examples per language, yielding 90 base examples and 360 judge responses.

121 **WMT MQM translation-quality contrast.** To test whether protocol sensitivity is equally
 122 severe for machine-translation evaluation, we sample 30 balanced examples per language
 123 pair from WMT MQM human-evaluation rows: en-de, en-ru, and zh-en. We use year-2022
 124 rows and derive high/low labels from within-pair MQM-score quantiles. This gives 90
 125 base examples. We run both reference-based judging (source, candidate, reference) and
 126 reference-free judging (source and candidate only). Each condition uses direct judging and
 127 explicit English-pivot judging for all pairs, plus target-language and bilingual rubrics for
 128 the two non-English target pairs, giving 600 WMT judge responses total.

129 **Stronger-judge experiment.** To check whether the main pattern is only a cheap-judge
 130 artifact, we run gpt-4.1-mini on a stratified candidate-quality subset: 25 examples per
 131 language/dimension, 200 base examples, and 800 judge responses. We also run a tar-
 132 getted WMT reference-free experiment on two cells where the main judge shows large or
 133 ambiguous protocol effects.

134 3.3 Judge Protocols

135 We compare four prompt surfaces. **P0 direct** presents the original candidate text with
 136 English instructions and an English rubric. **P1 target-rubric** presents the original candidate
 137 text with rubric wording translated into the target language when meaningful. **P2 explicit-**
 138 **pivot** translates the candidate, source, and/or reference material into English before judging,
 139 then asks the judge to score the English representation. **P3 bilingual** presents English and
 140 target-language rubric wording together while keeping the original candidate visible. The

Table 1: Dataset sampling summary for paper-facing item sets.

Run	Items	Langs	Dims	Pos	Neg	Source
candidate n50	400	4	2	200	200	0%
candidate experiment n25	200	4	2	96	104	0%
semantic n30	90	3	1	45	45	100%
wmt n30 shared items	90	3	1	45	45	100%
wmt ref-free experiment zh-en	30	1	1	15	15	100%
wmt ref-free experiment en-de	30	1	1	15	15	100%

summarization settings use all four protocols. The WMT setting uses P0 and P2 for all pairs and P1/P3 for non-English target pairs. All calls request strict JSON with a 1–5 score, coarse label, confidence, and a short rationale. A redacted prompt inventory is included in the release package so the system messages, rubric wording, protocol instructions, and JSON contracts can be audited without exposing item-specific text.

3.4 Metrics

We report both alignment and diagnostic metrics. Alignment is measured with Spearman correlation between judge score and human label, and AUROC for ranking human-positive examples above human-negative examples. Same-item protocol shift is the mean absolute score difference between two protocols on identical base examples:

$$\Delta_{p,q}^{\text{score}} = \frac{1}{n} \sum_{i=1}^n |s_{i,p} - s_{i,q}|. \quad (1)$$

For paired protocol comparisons, we report item-level AUROC deltas. We also report cross-language gaps, defined as the difference between the maximum and minimum per-language Spearman value under a fixed setting, post-hoc threshold calibration, JSON parse rate, returned token usage, and approximate API cost. Calibration uses small in-group calibration splits to choose a score threshold, then evaluates held-out balanced accuracy.

4 Results

4.1 Candidate-Quality Judging Is Highly Protocol-Sensitive

Table 2 summarizes the strongest candidate-quality findings. Protocol shifts are large on the same item, especially for target-language form dimensions. For Vietnamese grammar, target-language-rubric and English-pivot scores differ by 1.20 points on a 1–5 scale on average. For Vietnamese comprehensibility, bilingual and English-pivot scores differ by 0.96 points.

These shifts are not merely changes in score scale. In Turkish comprehensibility, direct judging reaches Spearman 0.60 and AUROC 0.84, while explicit English pivot drops to Spearman 0.33 and AUROC 0.69. The paired AUROC delta for pivot minus direct is -0.15. Thus a common workaround, translating into English before judging, can materially reduce alignment with human labels.

Qualitative inspection suggests two mechanisms. In one Vietnamese human-negative comprehensibility item, target-language judging scores the original summary 2, while the English pivot “Prepare materials” is scored 5 because the translation is clear but has removed the target-language error signal. Conversely, in a Turkish human-positive comprehensibility item, direct, target-language, and bilingual protocols all score 5, but the pivot translation repeats the sentence and the pivot judge assigns 1. Translation can therefore both clean away real errors and introduce new artifacts.

4.2 A Stronger Judge Does Not Remove the Failure Mode

The gpt-4.1-mini experiment is bounded, so we do not treat its exact correlations as final. It is useful as a stronger-model failure-mode check. On Turkish comprehensibility, the stronger

Table 2: Candidate-quality protocol sensitivity. Protocol columns report Spearman / AUROC; “Comp.” names the paired comparison for Δ AUROC.

Cell	Direct	Target	Pivot	Bil.	Comp.	Δ	Shift
tr comp.	0.60 / 0.84	0.56 / 0.81	0.33 / 0.69	0.57 / 0.82	Piv.-dir.	-0.15	0.62
vi comp.	0.35 / 0.69	0.16 / 0.59	0.09 / 0.55	0.40 / 0.71	Bil.-piv.	0.17	0.96
vi gram.	0.23 / 0.62	0.25 / 0.63	0.02 / 0.51	0.10 / 0.55	Piv.-tgt.	-0.12	1.20

judge gives direct judging Spearman 0.73 / AUROC 0.91 and bilingual judging 0.80 / 0.95, but explicit pivot falls to 0.17 / 0.59. The paired AUROC delta for pivot minus direct is -0.32, and bilingual minus pivot is +0.35. The same pattern appears for Vietnamese grammar: direct judging reaches 0.56 / 0.81, explicit pivot drops to -0.05 / 0.47, and bilingual judging recovers to 0.49 / 0.77.

Table 3: Stronger-judge experiment on a stratified n=25-per-cell subset.

Cell	Direct	Pivot	Bil.	Piv.-Dir. Δ	Bil.-Piv. Δ
tr comp.	0.73 / 0.91	0.17 / 0.59	0.80 / 0.95	-0.32	0.35
vi gram.	0.56 / 0.81	-0.05 / 0.47	0.49 / 0.77	-0.34	0.30

4.3 Modern Frontier Judges Preserve the Core Failure

We next run a focused modern-judge experiment on the highest-sensitivity cells from the gpt-4.1-mini experiment: Turkish comprehensibility, Vietnamese comprehensibility, Vietnamese grammar, and zh-en WMT reference-free translation quality. The experiment uses gpt-5.4-mini as a current low-cost judge and gpt-5.5 as a current frontier judge. It also adds one constructive mitigation, **P4 evidence-preserving bilingual judging**, which keeps the original target-language candidate visible and requires a short original-language evidence quote before scoring.

Table 4 shows that the English-pivot failure is not an artifact of the older mini judge. Under gpt-5.5, explicit pivot is still much worse than direct judging on Turkish comprehensibility (paired AUROC delta -0.26) and Vietnamese grammar (-0.47), and WMT zh-en reference-free judging also favors direct source-visible evaluation (-0.18). The mitigation is useful where pivot loses the relevant target-language evidence: P4 improves over explicit pivot by +0.27 AUROC on Turkish comprehensibility and +0.50 on Vietnamese grammar, while remaining competitive with direct and bilingual judging. This makes the paper’s practical recommendation sharper: for target-language form dimensions, keep original-language evidence visible and ask the judge to cite it.

Table 4: Modern-judge experiment and evidence-preserving mitigation. AUROC is computed within each n-per-protocol cell; deltas are paired AUROC differences. The WMT rows use only direct and pivot protocols.

Cell	Judge	Direct	Pivot	Bil.	P4	Piv.-Dir.	P4-Piv.
tr comp.	gpt-4.1-mini	0.91	0.59	0.95	–	-0.32	–
vi comp.	gpt-4.1-mini	0.79	0.62	0.75	–	-0.17	–
vi gram.	gpt-4.1-mini	0.81	0.47	0.77	–	-0.34	–
tr comp.	gpt-5.4-mini	0.95	0.65	0.88	–	-0.30	–
vi comp.	gpt-5.4-mini	0.78	0.68	0.81	–	-0.09	–
vi gram.	gpt-5.4-mini	0.84	0.44	0.82	–	-0.40	–
tr comp.	gpt-5.5	0.99	0.73	0.98	1.00	-0.26	0.27
vi comp.	gpt-5.5	0.81	0.69	0.84	0.83	-0.13	0.14
vi gram.	gpt-5.5	0.82	0.35	0.81	0.85	-0.47	0.50
WMT zh-en	gpt-4.1-mini	0.77	0.64	–	–	-0.14	–
WMT zh-en	gpt-5.4-mini	0.74	0.61	–	–	-0.13	–
WMT zh-en	gpt-5.5	0.80	0.62	–	–	-0.18	–

4.4 Semantic and MT Settings Clarify the Boundary

The source-grounded main-ideas setting gives a different but complementary pattern. Protocol shifts are smaller than for grammar/form dimensions, which is expected: semantic content can often survive translation better than grammar or fluency. Still, explicit pivot is not uniformly best. For Spanish main ideas, explicit pivot performs best. For Turkish, bilingual judging is best. For Vietnamese, all protocols are weak, and explicit pivot is chance-level. Paired comparisons reinforce this language dependence: bilingual minus pivot AUROC delta is +0.12 for Turkish and +0.11 for Vietnamese.

The WMT MQM contrast is intentionally smaller and narrower than the summarization runs, but it is important for claim boundaries. Reference-based translation-quality judging shows smaller direct-vs-pivot shifts. Reference-free MT judging surfaces stronger gpt-4o-mini protocol effects: for zh-en, direct judging reaches 0.60 / 0.82, while explicit pivot drops to 0.19 / 0.60. A targeted gpt-4.1-mini experiment preserves the zh-en direct-over-pivot direction but does not reproduce the en-de bilingual-over-pivot gain. Thus English pivoting is risky for target-language form and some source-grounded semantic settings, while the WMT evidence is best treated as protocol/model-sensitivity evidence rather than a stable global protocol ranking.

Table 5: Source-grounded XLSum main-ideas alignment by protocol. Values are Spearman / AUROC.

Lang.	Best	Best S/A	Direct	Pivot	Bil.
es-ES	P2 pivot	0.51 / 0.77	0.49 / 0.76	0.51 / 0.77	0.43 / 0.72
tr	P3 bil.	0.59 / 0.79	0.46 / 0.74	0.34 / 0.67	0.59 / 0.79
vi	P3 bil.	0.19 / 0.61	0.14 / 0.58	0.00 / 0.50	0.19 / 0.61

Table 6: WMT MQM translation-quality protocol extension. Values are Spearman / AUROC; protocol names are abbreviated.

Set.	Pair	Best	Best S/A	Direct	Target	Pivot	Bil.	Piv.-Dir.
Ref.	en-de	P3 bil.	0.47 / 0.69	0.39 / 0.67	0.32 / 0.60	0.40 / 0.68	0.47 / 0.69	0.00
Ref.	en-ru	P1 target	0.32 / 0.65	0.16 / 0.57	0.32 / 0.65	0.13 / 0.56	0.16 / 0.58	-0.01
Ref.	zh-en	P0 direct	0.35 / 0.69	0.35 / 0.69	n/a	0.30 / 0.67	n/a	-0.02
Free	en-de	P3 bil.	0.68 / 0.84	0.51 / 0.74	0.48 / 0.74	0.44 / 0.71	0.68 / 0.84	-0.03
Free	en-ru	P1 target	0.20 / 0.60	0.02 / 0.51	0.20 / 0.60	0.04 / 0.52	0.18 / 0.59	0.01
Free	zh-en	P0 direct	0.60 / 0.82	0.60 / 0.82	n/a	0.19 / 0.60	n/a	-0.22

4.5 No Single Protocol Dominates

Table 7 aggregates instability rather than declaring a universal winner. Across 17 main-experiment task/language cells, the best AUROC protocol is not direct in 11 cells, explicit pivot is the worst protocol in 10 cells, and 6 cells have at least one strongly supported paired AUROC change. The stronger experiments show the same qualitative point while tightening the WMT claim boundary: protocol choice is an experimental factor to report, not a nuisance prompt detail to hide.

Table 7: Protocol-instability summary by run. AUROC range is best minus worst protocol within a language/dimension cell.

Run	Cells	Med.	Max	$\geq .10$	Best $\neq$ dir.	Piv. worst	Sig.	Shift
candidate n50	8	0.07	0.17	3/8	4/8	4/8	2/8	1.20
semantic n30	3	0.11	0.12	2/3	3/3	2/3	1/3	0.70
wmt reference n30	3	0.08	0.09	0/3	2/3	2/3	0/3	0.53
wmt ref-free n30	3	0.13	0.22	2/3	2/3	2/3	3/3	0.70
cand. exp. n25	8	0.16	0.35	7/8	5/8	4/8	4/8	1.16
wmt free exp.	2	0.08	0.14	1/2	1/2	1/2	0/2	0.67

5 Analysis and Recommendations

Protocol choice changes what evidence is available. The strongest failures occur when the human label depends on target-language form. This matches the mechanism we expected: pivot translation can turn a target-language surface-form problem into fluent English, or introduce new English artifacts unrelated to the original candidate. For such dimensions, original-text judging with either direct or bilingual rubrics is safer than treating English pivot as a default.

Bilingual rubrics are useful but not magic. Bilingual judging often improves over explicit pivot in our highest-sensitivity cells, but it is not universally best. This suggests that multilingual evaluation should not replace one hidden default with another. Instead, papers should report multiple protocols on a small shared subset and state which protocol is used for the main score.

Aggregate scores hide language gaps. Protocol choice also changes disparity across languages. In candidate-quality comprehensibility, explicit pivot has a Spearman language gap of 0.63, from Vietnamese 0.09 to English 0.71. Bilingual judging reduces the same gap to 0.30. In source-grounded main ideas, explicit pivot again has the largest language gap, 0.51, from Vietnamese 0.00 to Spanish 0.51. A single aggregate judge score would hide these failures.

Calibration diagnoses score compression but does not solve protocol sensitivity. Post-hoc threshold calibration helps most in the source-grounded semantic setting, where judge scores are compressed low: with balanced human labels, the fixed score ≥ 4 threshold marks only 8.3% of examples as good, and the modal best in-group threshold is 2. In candidate-quality and reference-based WMT, calibration is neutral or slightly negative on average. Calibration is therefore useful as a score-threshold diagnostic, not a universal repair.

Repeatability is not the main explanation. The historical $n=20$ pilot and the $n=50$ main run overlap on 42 original-text prompt calls with identical prompt text. Across these repeated calls, exact score agreement is 92.9%, mean absolute score delta is 0.07, and the maximum score delta is 1. This is much smaller than the main protocol shifts, such as 0.96 for Vietnamese comprehensibility and 1.20 for Vietnamese grammar. By contrast, repeated explicit-pivot prompt IDs with regenerated translation text show lower agreement, suggesting that translation-pivot workflows add an additional source of pipeline volatility.

Global Judge Reporting Card. We recommend that multilingual LLM-as-a-judge papers report: judge model and version; prompt/rubric language; whether source, candidate, or reference texts were translated; whether target-language form dimensions were judged in the original language; per-language human alignment; same-item protocol-shift metrics; language gaps for each protocol; calibration split size and held-out effect; parse failure rate; returned token usage; and approximate cost per 1,000 judgments. Token usage is the stable accounting quantity; dollar values should be refreshed against current provider pricing before camera-ready release.

6 Limitations

The evidence is deliberately bounded. All main judge calls and all stronger/modern experiments use OpenAI-hosted models rather than an independent model family. Target-language rubric translations are pragmatic pilot translations, not native-speaker validated. The source-grounded setting covers one semantic dimension and three non-English languages. The WMT MQM contrast is small, covers one year and three language pairs, adds target/bilingual rubrics only for the two non-English target pairs, and uses quantile-derived binary labels from continuous MQM scores. The targeted WMT stronger-judge and modern-judge experiments are robustness checks rather than definitive model-family comparisons.

The reference-free WMT condition reuses the same MQM-derived labels, so it is a protocol stress test rather than a new human annotation target. SEAHORSE labels are binary yes/no labels for each dimension, while LLM judges produce 1–5 scores; thresholding and calibration are therefore part of the evaluation design.

7 Conclusion

Multilingual LLM-as-a-judge evaluation is under-specified without the judge protocol. The same item can receive different scores depending on whether the judge sees original text, translated text, target-language instructions, or bilingual instructions. English-pivot judging can help in some semantic and reference-based settings, but it can also erase target-language evidence, introduce translation artifacts, and increase language gaps. We recommend that global benchmark builders report protocol sensitivity as a first-class evaluation result.

References

- Andrea Bavaresco, Raffaella Bernardi, Leonardo Bertolazzi, Desmond Elliott, Raquel Fernández, Albert Gatt, Alexander Koller, Barbara Plank, David Schlangen, and Alberto Testoni. LLMs instead of human judges? a large scale empirical study across 20 NLP evaluation tasks. *arXiv preprint arXiv:2406.18403*, 2024.
- Elizabeth Clark, Shruti Rijhwani, Sebastian Gehrmann, Joshua Maynez, Roei Aharoni, Vitaly Nikolaev, Thibault Sellam, Aditya Siddhant, Dipanjan Das, and Ankur P. Parikh. SEAHORSE: A multilingual, multifaceted dataset for summarization evaluation. *arXiv preprint arXiv:2305.13194*, 2023.
- Alexander R. Fabbri, Wojciech Kryściński, Bryan McCann, Caiming Xiong, Richard Socher, and Dragomir Radev. SummEval: Re-evaluating summarization evaluation. *Transactions of the Association for Computational Linguistics*, 9:391–409, 2021.
- Markus Freitag, Ricardo Rei, Nitika Mathur, Chi-kiu Lo, Craig Stewart, George Foster, Alon Lavie, and Ondřej Bojar. Experts, errors, and context: A large-scale study of human evaluation for machine translation. *Transactions of the Association for Computational Linguistics*, 9:1460–1474, 2021.
- Markus Freitag, Ricardo Rei, Nitika Mathur, Chi-kiu Lo, Eleftherios Avramidis, Radu Rei, Brian Thompson, Tom Kocmi, Frederic Blain, Daniel Deutsch, Craig Stewart, Chrysoula Zerva, Sheila Castilho, and Alon Lavie. Results of WMT22 metrics shared task: Stop using BLEU – neural metrics are better and more robust. In *Proceedings of the Seventh Conference on Machine Translation*, 2022.
- Xiyan Fu and Wei Liu. How reliable is multilingual LLM-as-a-judge? *arXiv preprint arXiv:2505.12201*, 2025.
- Xiyan Fu and Wei Lu. When languages disagree: Self-evolving multilingual LLM judges. *arXiv preprint arXiv:2606.08092*, 2026.
- Tahmid Hasan, Abhik Bhattacharjee, Md. Saiful Islam, Kazi Mubasshir, Yuan-Fang Li, Yong-Bin Kang, M. Sohel Rahman, and Rifat Shahriyar. XL-sum: Large-scale multilingual abstractive summarization for 44 languages. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2021*, 2021.
- Junjie Hu, Sebastian Ruder, Aditya Siddhant, Graham Neubig, Orhan Firat, and Melvin Johnson. XTREME: A massively multilingual multi-task benchmark for evaluating cross-lingual generalization. In *Proceedings of the 37th International Conference on Machine Learning*, 2020.
- Tom Kocmi and Christian Federmann. Large language models are state-of-the-art evaluators of translation quality. In *Proceedings of the Eighth Conference on Machine Translation*, 2023.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG evaluation using GPT-4 with better human alignment. *arXiv preprint arXiv:2303.16634*, 2023.
- Joshua Maynez, Shashi Narayan, Bernd Bohnet, and Ryan McDonald. On faithfulness and factuality in abstractive summarization. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, 2020.
- Peiyi Wang, Lei Li, Liang Chen, Dawei Zhu, Binghui Lin, Yunbo Cao, and Zhifang Sui. Large language models are not fair evaluators. *arXiv preprint arXiv:2305.17926*, 2023.
- Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhenghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric P. Xing, Hao Zhang, Joseph E. Gonzalez, and Ion Stoica. Judging LLM-as-a-judge with MT-bench and chatbot arena. In *Advances in Neural Information Processing Systems*, 2023.